

Adversarial Search

LESSON 6

Reading

Chapter 6

Outline

- Last time: Heuristic, Informed search
 - $h(x)$: utility function
- Optimal decisions
- α - β pruning
- Imperfect, real-time decisions

Games vs. Search problems

- “Unpredictable” opponent -> specifying a move for every possible opponent reply
- Time limits -> unlikely to find goal, must Approximate
- Hmm: Is Ataxx a game or a search problem by this definition?


Two player games

Max always moves first.

Min is the opponent.

We have


- An initial state.
- A set of operators.
- A terminal test (which tells us when the game is over).
- A utility function (evaluation function).



Max Vs Min

The utility function is like the heuristic function we have seen in the past, except it evaluates a node in terms of how good it is for each player. Positive values indicate states advantageous for Max, negative values indicate states advantageous for Min.

Game tree (2-player, deterministic, turns)




Key property: we have a **zero-sum** game.

- Loosely, it means that there's a loser for every winner.
- Total utility score over all agents sum to zero.
- Makes the game adversarial.



To think about: not zero sum?


A simple abstract game.
Max makes a move, then **Min** replies.



An action by one player is called a *ply*, two ply (a action and a counter action) is called a *move*.


The Minimax Algorithm

- Generate the game tree down to the terminal nodes.
- Apply the utility function to the terminal nodes.
- For a **S** set of sibling nodes, pass up to the parent...
 - the lowest value in **S** if the siblings are 
 - the largest value in **S** if the siblings are 
- Recursively do the above, until the backed-up values reach the initial state.
- The value of the initial state is the minimum score for Max.




In this game Max's best move is A1, because he is guaranteed a score of at least 3.


Example: Dodgen game



Example: Dodgen game



Example: Dodgen game



Evaluation Function

```
function MaxVal(u);  
begin  
    if u là đỉnh kết thúc then MaxVal(u)  $\leftarrow f(u)$   
    else MaxVal(u)  $\leftarrow \max\{MinVal(v) \mid v \text{ là đỉnh con của } u\}$   
end;
```


```
function MinVal(u);
```

?


Evaluation Function

■ Eval(u)

- $\text{Eval}(u) > 0$: prefer Max
- $\text{Eval}(v) < 0$: prefer Min



?




?

30	35	40
15	20	25
0	5	10


Giá trị quân Trắng.

-10	-25	-40
-5	-20	-35
0	-15	-30

Giá trị quân Đen.




Trắng cản trực tiếp Đen
được thêm 40 điểm.




Trắng cản gián tiếp Đen
được thêm 30 điểm.

Evaluation Function



75



-5

Minimax Algorithm

```
function MINIMAX-DECISION(state) returns an action  
   $v \leftarrow \text{MAX-VALUE}(\textit{state})$   
  return the action in SUCCESSORS(state) with value v
```

```
function MAX-VALUE(state) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow -\infty$   
  for a, s in SUCCESSORS(state) do  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s))$   
  return v
```

```
function MIN-VALUE(state) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow \infty$   
  for a, s in SUCCESSORS(state) do  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s))$   
  return v
```


Properties of Minimax

- **Complete?** Yes (if tree is finite)
- **Optimal?** Yes (against an optimal opponent)
- **Time complexity?** $O(b^m)$, where b is the effective branching factor and m is the depth of the terminal states.
- **Space complexity?** $O(b^m)$ (depth-first exploration)


For chess, $b \approx 35$, $m \approx 100$ for “reasonable” games

-> exact solution completely infeasible

What can we do?


PRUNING!

α - β Pruning




- Pruning *sub-tree a* if $Eval(u) > Eval(v)$


α - β Pruning example




α - β Pruning example




α - β Pruning example



α - β Pruning example




α - β Pruning example



Why is it called α - β

- α is the value of the best (i.e., highest-value) choice found so far at any choice point along the path for *max*
- If v is worse than α , *max* will avoid it
 - prune that branch
- Define β similarly for *min*



The α - β Algorithm

function ALPHA-BETA-SEARCH(*state*) *returns an action*

inputs: *state*, current state in game

$v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$

return the *action* in SUCCESSORS(*state*) with value v

function MAX-VALUE(*state*, α , β) *returns a utility value*

inputs: *state*, current state in game

α , the value of the best alternative for MAX along the path to *state*

β , the value of the best alternative for MIN along the path to *state*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow -\infty$

for a, s in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s, \alpha, \beta))$

if $v \geq \beta$ **then return** v

$\alpha \leftarrow \text{MAX}(\alpha, v)$


return v

The α - β Algorithm

```
function MIN-VALUE( $state, \alpha, \beta$ ) returns a utility value
  inputs:  $state$ , current state in game
            $\alpha$ , the value of the best alternative for MAX along the path to  $state$ 
            $\beta$ , the value of the best alternative for MIN along the path to  $state$ 

  if TERMINAL-TEST( $state$ ) then return UTILITY( $state$ )
   $v \leftarrow +\infty$ 
  for  $a, s$  in SUCCESSORS( $state$ ) do
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s, \alpha, \beta))$ 
    if  $v \leq \alpha$  then return  $v$ 
     $\beta \leftarrow \text{MIN}(\beta, v)$ 
  return  $v$ 
```

The α - β Algorithm: example



Resource Limits

- The big problem is that the search space in typical games is very large.
- Suppose we have 100 secs, explore 104 nodes/sec -> 106 nodes per move

Standard approach:

- cutoff test:
 - e.g., depth limit
- evaluation function
 - = estimated desirability of position

Evaluation functions

- For chess, typically **linear** weighted sum of **features**
$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$
- e.g., $w_1 = 9$ with
 $f_1(s) = (\# \text{ of white queens}) - (\# \text{ of black queens}), \text{ etc.}$
- Caveat: assumes independence of the features
 - Bishops in Western chess better at endgame
 - Unmoved king and rook needed for castling
- Should model the *expected utility value* states with the same feature values lead to.

Other problems

- Applying utility functions on end game scenarios may not solve the game
- Use a policy or lookup table (taken from previous game history)
- Stochastic games
- Calculate the expected value of a position

What do you need to do

- Implement Minimax
- Implement Pruning (optional)
- Implement an evaluation function
 - ✓ Input: board, selected grid location
 - ✓ Output: continuous value
- (really optional) use state

Summary

- ❑ Games are fun to work on!
- ❑ They illustrate several important points about AI
- ❑ Perfection is unattainable -> must approximate
- ❑ Good idea to think about what to think about